

# Ibrahim Yaghi

AI Platform Engineer & Agentic Systems Architect

ibrahem.yaghi@gmail.com · +962 78 503 1314 · linkedin.com/in/ibrahim-yaghi · Amman, Jordan

~4M JOD

Cost Savings Delivered

■ Best Digital Transformation Initiative in Telecom — Global Economics 2025

I build AI systems that actually ship — autonomous agents, intelligent pipelines, and developer platforms running in production at Zain Jordan, one of the region's leading telecoms. Below is a portfolio of the AI projects I designed, engineered, and delivered.

## Agentic AI Platform — Zain Jordan

01

### Agent Orchestrator (In-House)

Custom orchestration engine built from scratch to manage all internal agents and Python backends

- Designed and built a proprietary orchestration layer from scratch — no third-party frameworks — to route user requests across multiple specialized agents.
- Unified interface for all internal AI services, connecting Python backends, Redis message queues, and a lightweight SQLite state store.
- Handles routing logic, session context, agent chaining, fallback handling, and response aggregation in real time.
- Enabled rapid deployment of new agents by providing a standardized integration contract for all services.

**Impact:** Became the backbone of Zain's internal AI ecosystem, enabling all agents below to operate under a single, coherent platform.

Python

Vanilla JS

HTML

Redis

SQLite

Custom  
Architecture

02

### Agentic AI — Sales Data Agent

Autonomous agent converting natural-language questions into executed SQL queries and visual charts

- Sales employees chat with the agent in plain language; the agent autonomously analyzes intent, generates parameterized PL/SQL queries, and executes them securely.
- Returns results as charts and formatted data — no database access or SQL knowledge needed from end users.
- Built with strict no-data-exposure controls: the agent never returns raw database dumps.
- Eliminated manual reporting cycles across multiple business departments.

**Impact:** Replaced full manual reporting cycles; adopted across sales, finance, and operations teams.

Python

LangChain

OpenAI

PL/SQL

Data Visualization

03

**Marketing Intelligence Agent***Retrieval agent that monitors competitor offers and generates actionable recommendations*

- Autonomously reads competitor websites to extract current offers, devices, and packages.
- Compares retrieved competitor data against Zain's live offerings side-by-side.
- Generates strategic suggestions for new packages, pricing changes, or promotional offers.
- Enables the marketing team to stay competitive without manual research overhead.

Python

LangChain

OpenAI

Web Retrieval

RAG

04

**LinkedIn Talent Hunting Agent***HR-facing conversational agent that searches LinkedIn for candidates by skill and education*

- HR staff describe the candidate profile they need in natural language — skills, education, experience level.
- Agent searches LinkedIn and retrieves the top 10 best-matching profiles automatically.
- Reduces sourcing time drastically compared to manual LinkedIn searches.
- Supports complex queries combining multiple skill sets and qualifications.

Python

LangChain

OpenAI

LinkedIn Data

NLP

05

**Interview Screening Agent***AI phone screening agent that scores CVs and conducts structured voice interviews*

- Before the interview, the agent scores the candidate's CV against the job requirements and auto-generates tailored interview questions.
- Conducts a conversational phone screening in English with the candidate.
- Delivers a structured evaluation report to the HR team, accelerating the shortlisting process.
- Reduces recruiter time spent on early-stage screening rounds.

Python

LangChain

OpenAI

NLP

Voice AI

06

**Key Account Manager (KAM) Agent***Action-capable agent for KAMs to manage and query mobile number operations*

- KAMs interact conversationally to perform operational tasks: check number availability, reassign mobile numbers, query account statuses.
- Agent executes real actions against backend telecom systems via tool calling — not just retrieval.
- Reduces dependency on manual back-office processes and ticketing systems.

Python

LangChain

OpenAI

Telecom APIs

Tool Calling

07

**CRM Testing Agent***Conversational agent for testing the in-house CRM through natural-language test scenarios*

- QA teams describe test scenarios in plain language; the agent executes them against the new CRM.
- Accelerates testing cycles for new CRM features without requiring manual test scripting.
- Captures and reports test outcomes in structured format for QA review.

Python

LangChain

OpenAI

QA Automation

CRM

08

**Technical Documentation Generator Agent***Agent that auto-generates technical documentation at the end of a development cycle*

- Takes project code, architecture notes, and development context as input.
- Produces structured technical documentation ready for review — reducing the documentation burden on engineers.
- Standardizes documentation format across all internal AI and software projects.

Python

LangChain

OpenAI

NLP

Documentation

**Knowledge & RAG Platform**

09

**Botzilla — Enterprise Knowledge Chatbot***Multi-team RAG chatbot with strict document isolation per team, powered by Qdrant*

- Each team operates in a fully isolated knowledge namespace — a team can only retrieve documents belonging to their own corpus, enforced at the vector DB layer.
- Built on Qdrant as the vector database for fast, scalable semantic search over private document collections.
- Returns accurate, grounded answers in 5–15 seconds over private enterprise document corpora.
- Runs entirely within the enterprise-controlled environment — no data leaves the organization.
- Fully customizable to new teams and document sets without re-engineering the core platform.

**Impact:** Deployed across multiple departments, replacing manual document search and enabling self-serve knowledge retrieval.

Python

LangChain

OpenAI

Qdrant

RAG

Vector DB

**Platform Engineering & Infrastructure**

10

**Reconciliation Pipeline Optimization***PySpark-based rewrite of enterprise reconciliation pipelines — 8 hours → minutes*

- Identified critical performance bottlenecks in large-scale enterprise reconciliation jobs.
- Re-architected batch pipelines using PySpark distributed processing, reducing runtime from 8 hours to under 10 minutes.
- Improved downstream data reliability for AI workloads and financial reporting systems.

**Impact:** Direct contributor to the ~4M JOD (~5.6M USD) in cumulative cost savings delivered at Zain Jordan.

Python

PySpark

Distributed Computing

Data Engineering